

Unsupervised Learning

CS446-Machine Learning

CS444-Deep Learning for Computer Vision

Qi Long

Contents

0	Preliminaries	4
0.1	Lagrangian Multiplier	4
0.1.1	Goal: Find Optimum	4
0.1.2	Method: same derivative	4
0.2	Linear Algebra Basics	4
0.2.1	Inverse	5
0.2.2	Trace	5
0.3	Matrix Derivative	5
0.3.1	First Order	5
0.3.2	Second Order	5
0.4	Norm	5
0.4.1	Vector Norm	5
0.4.2	Matrix Norm	6
0.5	SVD	6
0.5.1	Unitary Matrix	6
0.5.2	Singular Value Decomposition	6
0.6	Gaussian Distribution	7
0.7	Covariance	7
0.8	Conditional Probability	7
0.8.1	Markov Chain	7
0.8.2	Bayes Theorem	7
0.9	Information Theory	7
0.9.1	Shannon Entropy	7
0.9.2	Cross Entropy	8
0.9.3	KL-Divergence	8
0.9.4	TV-Distance	9

0.9.5	Pinsker's Inequality	9
0.9.6	Mutual Information	9
0.10	Jensen's Inequality	9
1	PCA	9
1.1	Goal: Dimension Reduction	9
1.2	Preprocessing: Center Data	9
1.3	Derivation 1: Direction Maximize Variance	10
1.3.1	Objective	10
1.3.2	Lagrangian	10
1.4	Algorithm 1	10
1.4.1	Find Lower-dim Directions	10
1.4.2	Compress Data	10
1.4.3	Reconstruct Data	10
1.5	Derivation 2: Direction Minimize Reconstruction Loss	11
1.5.1	Objective	11
1.6	Solve by SVD	11
2	Clustering	12
2.1	K-Means	12
2.1.1	Objective	12
2.1.2	Alternate Optimization	12
2.1.3	K-Means++	12
2.2	Gaussian Mixture Models	13
2.2.1	Single Gaussian	13
2.2.2	Mixed Gaussian	13
2.3	Evaluation	14
3	VAE	14
3.1	Setup	14
3.2	ELBO	15
3.2.1	From Objective to ELBO	15
3.2.2	From ELBO to Loss	15
3.2.3	Methodology	16
3.3	Training	17
4	GAN	17
4.1	Binary Classification Game	17
4.2	Non-saturating GAN Loss (NSGAN)	18
4.3	Optimization	19
4.4	Advanced Architectures	20

4.4.1	DCGAN	20
4.4.2	WGAN	20
4.4.3	LSGAN	20
4.4.4	ProgressiveGAN	20
4.4.5	StyleGAN	20
4.4.6	BigGAN	22
4.4.7	StyleGAN-XL	22
4.4.8	GigaGAN	22
4.5	Evaluation	23
4.5.1	Human Studies	23
4.5.2	Inception Score (IS)	23
4.5.3	Frechet Inception Distance (FID)	23
5	Diffusion	23
5.1	Setup	23
5.2	Encoder	24
5.2.1	Objective	24
5.2.2	Derivation	24
5.3	ELBO	25
5.4	Disastrous Loss Function	26
5.5	Reverse Encoder	26
5.6	Loss Function	27
5.7	Training	29
5.7.1	Consistency Loss	29
5.7.2	Reconstruction Loss	29
5.7.3	Training Algorithm	30
5.8	Inference	30
5.9	Advanced Architectures	30
5.9.1	DALL-E 2	30
5.9.2	Latent Diffusion	31
5.9.3	Google Imagen	31
5.9.4	SDXL	31
5.9.5	Progressive Distillation	31
5.9.6	Latent Consistency Models	32

0 Preliminaries


0.1 Lagrangian Multiplier

0.1.1 Goal: Find Optimum

Find the optimum of a function given a restriction of another function, for example,

$$\max_{g(x,y)=c} f(x,y)$$

0.1.2 Method: same derivative



$f(x,y) = d_n$ can be regarded as contours with **adjustable** values d_n .

$g(x,y) = c$ can be regarded as a fixed curve crossing contours.

Then the optimum is when $f(x,y) = d_n$ and $g(x,y) = c$ have same derivative.

Import λ to represent adjustable d_n ,

$$\nabla \frac{1}{\lambda} f(x,y) = \nabla (g(x,y) - c)$$

$$\nabla [f(x,y) - \lambda(g(x,y) - c)] = 0$$

Solve this gives λ , plug back to objective to solve optimum.

0.2 Linear Algebra Basics

$$A(B + C) = AB + AC$$

$$(A + B)C = AC + BC$$

0.2.1 Inverse

$$(ABC \dots)^{-1} = \dots C^{-1} B^{-1} A^{-1}$$

$$(A^T)^{-1} = (A^{-1})^T$$

$$(A + B)^T = A^T + B^T$$

$$(AB)^T = B^T A^T$$

0.2.2 Trace

$$\text{Tr}(ABC) = \text{Tr}(BCA) = \text{Tr}(CAB)$$

$$\text{Tr}(A + B) = \text{Tr}(A) + \text{Tr}(B)$$

0.3 Matrix Derivative

0.3.1 First Order

$$\frac{\partial x^T a}{\partial x} = \frac{\partial a^T x}{\partial x} = a$$

$$\frac{\partial a^T X b}{\partial X} = ab^T$$

0.3.2 Second Order

$$\frac{\partial x^T B x}{\partial x} = (B + B^T)x$$

$$\frac{\partial b^T X^T X c}{\partial X} = X(bc^T + cb^T)$$

0.4 Norm

0.4.1 Vector Norm

General **p-norm**:

$$\|x\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}, p \geq 1$$

2-norm or **Euclidean norm**: Euclidean distance,

$$\|x\|_2 = \sqrt{\sum_{i=1}^n x_i^2}$$
$$\|x\|_2^2 = x^T x$$

1-norm: Manhattan distance, grid path length,

$$\|x\|_1 = \sum_{i=1}^n |x_i|$$

max-norm: Max-dim distance,

$$\|x\|_\infty = \max_i |x_i|$$

0.4.2 Matrix Norm

Frobenius Norm: sum of squared each entry,

$$\|A\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n a_{ij}^2} = \sqrt{\text{tr}(A^T A)}$$

0.5 SVD

For matrix M of size $(m \times n)$,

$$M = U \Sigma V^T$$

0.5.1 Unitary Matrix

$$U^{-1} = U^T$$

$$U^* U = U U^* = I$$

0.5.2 Singular Value Decomposition

- U : size of $(m \times m)$, unitary matrix.
- Σ : size of $(m \times n)$, diagonal matrix uniquely defined by M , diagonal entries are **singular values** of M .
- V : size of $(n \times n)$, unitary matrix.

0.6 Gaussian Distribution

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

0.7 Covariance

$$\text{Cov}[X, Y] = \mathbb{E}[(X - \mathbb{E}(X))(Y - \mathbb{E}(Y))] = \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y]$$

$$\Sigma = \begin{pmatrix} \text{Var}[X_1] & \text{Cov}[X_1, X_2] & \cdots & \text{Cov}[X_1, X_D] \\ \text{Cov}[X_2, X_1] & \text{Var}[X_2] & \cdots & \text{Cov}[X_2, X_D] \\ \vdots & \vdots & \ddots & \vdots \\ \text{Cov}[X_D, X_1] & \text{Cov}[X_D, X_2] & \cdots & \text{Var}[X_D] \end{pmatrix}$$

0.8 Conditional Probability

0.8.1 Markov Chain

If X, Y, Z form a Markov Chain ($X \rightarrow Y \rightarrow Z$), then

$$p(x, y, z) = p(x)p(y|x)p(z|y)$$

0.8.2 Bayes Theorem

$$p(x|y, z) = p(y|x, z) \frac{p(x|z)}{p(y|z)}$$


0.9 Information Theory

0.9.1 Shannon Entropy

Higher value corresponds to more uncertainty.

$$H(p) = - \sum_{i=1}^n p_i \log p_i$$

$$H(X) = - \sum_{k=1}^K p(X = k) \log p(X = k) = -\mathbb{E}[\log p(X)]$$



0.9.2 Cross Entropy

General case:

$$H_{ce}(p, q) = - \sum_{k=1}^K p_k \log q_k$$

Binary case:

$$H_{ce}(p, q) = -p \log q - (1 - p) \log(1 - q)$$

0.9.3 KL-Divergence

Non-negative Evaluation of similarity between two distributions. Low value corresponds to similar distribution.

$$D_{KL}(P||Q) = \sum_{x \in X} P(x) \log \frac{P(x)}{Q(x)}$$

Bernoulli Distributions:

$$D_{KL}(P||Q) = p \log \frac{p}{q} + (1 - p) \log \frac{1 - p}{1 - q}$$

Gaussian Distributions:

$$D_{KL}(P||Q) = \frac{1}{2}(\mu_1 - \mu_0)^T \Sigma^{-1}(\mu_1 - \mu_0) = \frac{1}{2} \|\mu_0 - \mu_1\|_{\Sigma}^2$$

*The KL-Divergence of two **identical-variance** Gaussians is just Euclidean distance square between two mean vectors.

0.9.4 TV-Distance

An **event** (contain many x) that has largest difference between distribution P and Q.

$$d_{TV}(P, Q) = \sup_{A \subseteq \mathbb{R}^d} |P(A) - Q(A)|$$

0.9.5 Pinsker's Inequality

$$d_{TV}(\cdot, \cdot) \leq \sqrt{\frac{1}{2} D_{KL}(\cdot || \cdot)}$$

0.9.6 Mutual Information

Reduction in uncertainty given another variable

$$I(X; Y) = H(X) - H(X|Y) = H(Y) - H(Y|X) = D_{KL}(P(X, Y) || P(X)P(Y))$$

0.10 Jensen's Inequality

For any random variable X and any **concave** function f , it holds that

$$f(\mathbb{E}[X]) \geq \mathbb{E}[f(X)]$$

1 PCA

1.1 Goal: Dimension Reduction

Find a **small number of “directions”** (low-dim linear subspace) in input space that explain **variation** in input data; re-represent data by projecting along those directions.

1.2 Preprocessing: Center Data

In order to let variance (1) get rid of mean,

$$\mu = \frac{1}{N} \sum_i x_i$$

$$\bar{X} = X - \mu$$

where $\sum_i x_i$ is sum of each dimension over all data points.

$X \rightarrow$ each data point each dim – mean of that dim over all data points.

1.3 Derivation 1: Direction Maximize Variance

1.3.1 Objective

$$\max_{\|w\|_2^2=1} \text{Var}(w^T \bar{x}) = \max_{\|w\|_2^2=1} \mathbb{E}[(w^T \bar{x} - 0)(\bar{x}^T w - 0)] = \max_{\|w\|_2^2=1} w^T \Sigma w \quad (1)$$

where Σ is covariance matrix,

$$\Sigma = \frac{1}{N}(\bar{x} - 0)(\bar{x}^T - 0) = \frac{1}{N}\bar{x}\bar{x}^T$$

1.3.2 Lagrangian

$f(x) = w^T \Sigma w$, $g(x) = \|w\|_2^2 - 1 = w^T w$, get Lagrangian

$$L(w, \lambda) = w^T \Sigma w - \lambda(w^T w - 1)$$

$$\frac{\partial L(w, \lambda)}{\partial w} = (\Sigma + \Sigma^T)w - \lambda(2w) = 0$$

$$\Sigma w = \lambda w$$

Therefore, λ is eigenvalue of Σ . Plug this back to objective get

$$\max w^T \Sigma w = \max \lambda w^T w = \max \lambda \|w\|_2^2 = \max \lambda$$

Therefore, λ is largest eigenvalue of Σ , w is corresponding eigenvector.

1.4 Algorithm 1

1.4.1 Find Lower-dim Directions

Take n highest eigenvalues and eigenvectors $U = [w_1, w_2 \dots]$.

1.4.2 Compress Data

Project data to linear subspace

$$\hat{x} = U^T(x - \mu) \quad (2)$$

1.4.3 Reconstruct Data

Back projection

$$\tilde{x} = U\hat{x} + \mu \quad (3)$$

1.5 Derivation 2: Direction Minimize Reconstruction Loss

1.5.1 Objective

$$L(W) = \frac{1}{N} \sum_{n=1}^N \|x_n - \text{reconstruct}(\text{compress}(x_n, w))\|_2^2$$

Using (2), (3),

$$\begin{aligned} L(W) &= \frac{1}{N} \sum_{i=1}^N \|x^{(i)} - ww^T x^{(i)}\|_2^2 = \frac{1}{N} \|\bar{X} - ww^T \bar{X}\|_F^2 \\ &= \frac{1}{N} \text{Tr}(((I - ww^T)\bar{X})^T((I - ww^T)\bar{X})) = \frac{1}{N} \text{Tr}(\bar{X}\bar{X}^T(I - ww^T)^T(I - ww^T)) \end{aligned}$$

Since projection $I - ww^T$ has the property

$$(I - ww^T)^T(I - ww^T) = I - 2ww^T + w(w^T w)w^T = I - 2ww^T + \|w\|_2^2 ww^T = I - ww^T$$

$$L(W) = \frac{1}{N} \text{Tr}(\bar{X}\bar{X}^T(I - ww^T)^T(I - ww^T)) = \text{Tr}(\Sigma(I - ww^T)) = \text{Tr}(\Sigma) - \text{Tr}(\Sigma ww^T)$$

because $\text{Tr}(\Sigma)$ is fixed wrt w , objective becomes

$$\max_{\|w\|_2^2=1} w^T \Sigma w$$

Then the rest of derivation is the same as Derivation 1-Lagrangian.

1.6 Solve by SVD

To compute

$$\Sigma = \frac{1}{N} \bar{X}\bar{X}^T$$

let

$$\frac{1}{\sqrt{N}} \bar{X} = USV^T$$

then

$$\Sigma U = USV^T V S U^T U = S^2 U$$

Therefore, need to compute Singular Values S and U .

2 Clustering

2.1 K-Means

2.1.1 Objective

Assign each point to a cluster, minimize the total Euclidean distance of data points to clusters' means.

$$\min_{\mu} \min_r \sum_{i \in D} \sum_{k=1}^K \frac{1}{2} r_{ik} \|x^{(i)} - \mu_k\|_2^2$$

where D is dataset, K is number of clusters, r_i is one-hot ($r_{ik} \in 0, 1$, $\sum_{k=1}^K r_{ik} = 1$).

2.1.2 Alternate Optimization

Given μ fixed, update r (assign each point to nearest cluster mean, $O(KNd)$),

$$r_{ik} = \begin{cases} 1, & k = \operatorname{argmin}_{k \in 1, \dots, K} \|x^{(i)} - \mu_k\|_2^2 \\ 0, & \text{otherwise} \end{cases}$$

Given r fixed, update μ (recalculate mean of each cluster, $O(Nd)$),

$$\nabla_{\mu_k} L = \sum_{i \in D} r_{ik} (x^{(i)} - \mu_k) = 0$$

$$\mu_k = \frac{\sum_{i \in D} r_{ik} x^{(i)}}{\sum_{i \in D} r_{ik}}$$

Iterate until convergence (no more update).

2.1.3 K-Means++

Initialization: randomly choose first center, pick new centers with

$$p \propto \|x^{(i)} - \mu_k\|_2^2$$

(farthest from previous centers).

2.2 Gaussian Mixture Models

2.2.1 Single Gaussian

Goal: use a Gaussian to approximate the data distribution.

By tuning parameters μ, σ , maximize **generative** objective:

$$p(x^{(i)}|\mu, \sigma) = N(x^{(i)}|\mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(x^{(i)} - \mu)^2\right)$$

minimize negative log likelihood:

$$L = -\log \prod_{i \in D} p(x^{(i)}|\mu, \sigma) = \frac{N}{2} \log(2\pi\sigma^2) + \sum_{i \in D} \frac{1}{2\sigma^2} (x^{(i)} - \mu)^2$$

$$\frac{\partial L}{\partial \mu} = \sum_{i \in D} \frac{-1}{\sigma^2} (x^{(i)} - \mu) = 0 \quad \rightarrow \quad \mu^* = \frac{1}{N} \sum_{i \in D} x^{(i)}$$

$$\frac{\partial L}{\partial \sigma} = \sum_{i \in D} \frac{N}{2} \cdot \frac{4\pi\sigma}{2\pi\sigma^2} - \frac{1}{\sigma^3} (x^{(i)} - \mu)^2 = 0 \quad \rightarrow \quad \sigma^{2*} = \frac{1}{N} \sum_{i \in D} (x^{(i)} - \mu)^2$$

2.2.2 Mixed Gaussian

Soft version of K-Means:

- Each sample is partially assigned to **all** clusters (with responsibility r_{ik} , weight π_k is shared among samples).

$$p(x^{(i)}|\pi, \mu, \sigma) = \sum_{k=1}^K \pi_k N(x^{(i)}|\mu_k, \sigma_k)$$

$$r_{ik} = \frac{\pi_k N(x^{(i)}|\mu_k, \sigma_k)}{\sum_{\hat{k}=1}^K \pi_{\hat{k}} N(x^{(i)}|\mu_{\hat{k}}, \sigma_{\hat{k}})}$$

- Each cluster's mean is updated biased wrt to samples. (Each cluster is a Gaussian that only takes partial consideration (r_{ik}) of each sample)

$$N_k = \sum_{i \in D} r_{ik}$$

Minimize objective by adjusting parameters π, μ, σ :

$$L = -\log \prod_{i \in D} p(x^{(i)} | \pi, \mu, \sigma) = -\sum_{i \in D} \log \sum_{k=1}^K \pi_k N(x^{(i)} | \mu_k, \sigma_k)$$

where π_k is the assigned partial of sample to cluster k , $\sum_{k=1}^K \pi_k = 1$.
Similar to single Gaussian except small modifications,

$$\frac{\partial L}{\partial \mu_k} = \sum_{i \in D} r_{ik} \frac{-1}{\sigma^2} (x^{(i)} - \mu) = 0 \quad \rightarrow \quad \mu_k^* = \frac{1}{N_k} \sum_{i \in D} r_{ik} x^{(i)}$$

$$\frac{\partial L}{\partial \sigma_k} = \sum_{i \in D} r_{ik} \left(\frac{N}{2} \cdot \frac{4\pi\sigma}{2\pi\sigma^2} - \frac{1}{\sigma^3} (x^{(i)} - \mu)^2 \right) = 0 \quad \rightarrow \quad \sigma_k^{2*} = \frac{1}{N_k} \sum_{i \in D} r_{ik} (x^{(i)} - \mu)^2$$

For updating π_k , use Lagrangian,

$$\begin{aligned} \frac{\partial L}{\partial \pi_k} &= \sum_{i \in D} \frac{N(x^{(i)} | \mu_k, \sigma_k)}{\sum_{\hat{k}=1}^K \pi_{\hat{k}} N(x^{(i)} | \mu_{\hat{k}}, \sigma_{\hat{k}})} + \lambda = 0 \\ \sum_{i \in D} \frac{\sum_{\hat{k}=1}^K \pi_{\hat{k}} N(x^{(i)} | \mu_{\hat{k}}, \sigma_{\hat{k}})}{\sum_{\hat{k}=1}^K \pi_{\hat{k}} N(x^{(i)} | \mu_{\hat{k}}, \sigma_{\hat{k}})} + \sum_{k=1}^K \pi_k \lambda &= 0 \\ N + \lambda &= 0 \quad \rightarrow \quad \lambda = -N, \pi_k = \frac{N_k}{N} \end{aligned}$$

2.3 Evaluation


- Reconstruction Loss.
- Purity (if label available).

3 VAE

3.1 Setup

- Encoder: deep net to predict mean and var based on input $x^{(i)}$, so for a different $x^{(i)}$ will get a different Gaussian.

$$q_\phi(z|x) = N(z; \mu_\phi(x), \sigma_\phi(x))$$



- Sample latent var: parameterization trick

$$z \sim q_\phi(z|x) = N(z; \mu_\phi(x), \sigma_\phi(x)) = \mu_\phi(x) + \sigma_\phi(x) \cdot \epsilon$$

where $\epsilon \sim N(0, 1)$

- Decoder: deep net to predict x from latent var. During Inference time, only decoder is used (with a sample from Normal Distribution as input).

$$\hat{x} = p_\theta(x|z)$$

3.2 ELBO

3.2.1 From Objective to ELBO

Objective is to maximize the probability of decoder generating ground truth x .
Marginalize all latent var z :

$$\log p_\theta(x) = \log \int p_\theta(x, z) dz = \log \int q_\phi(z|x) \frac{p_\theta(x, z)}{q_\phi(z|x)} dz$$

By Jensen-inequality:

$$\geq \int q_\phi(z|x) \log \frac{p_\theta(x, z)}{q_\phi(z|x)} dz = L(p_\theta, q_\phi)$$

3.2.2 From ELBO to Loss

The term $p_\theta(x, z)$ is intractable, use bayes to separate it into two tractable losses:

$$\begin{aligned} L(p_\theta, q_\phi) &= \int q_\phi(z|x) \log \frac{p_\theta(x, z)}{q_\phi(z|x)} dz = \int q_\phi(z|x) \log \frac{p_\theta(x|z)p(z)}{q_\phi(z|x)} dz \\ &= \int q_\phi(z|x) \log \frac{p(z)}{q_\phi(z|x)} dz + \int q_\phi(z|x) \log p_\theta(x|z) dz = -D_{KL}(q_\phi, p) + \mathbb{E}[\log p_\theta(x|z)] \end{aligned}$$

Interpretation of loss:

- Prior Matching (Regularization): restrict $q_\phi(z|x)$ close to $N(z; 0, 1)$ ($p(z)$) to avoid encoder directly copying input instead of discovering features.

$$-D_{KL}(q_\phi||p) = \int q_\phi(z|x) \log \frac{p(z)}{q_\phi(z|x)} dz$$

- Reconstruction: minimize $(x - \text{decode}(\text{encode}(x)))$.

$$\mathbb{E}[\log p_\theta(x|z)] = \int q_\phi(z|x) \log p_\theta(x|z) dz$$

3.2.3 Methodology

From Objective to ELBO: transform $p_\theta(x)$ to condition on $q_\phi(z|x)$.

- Goal: log probability of predicting real data. (**Fixed**)

$$\log p_\theta(x_0)$$

- Marginalize all hidden / latent variables.

$$\log p_\theta(x_0) = \log \int p_\theta(x_0, z_{1:T}) dz$$

- Choose q_ϕ to represent **whole** forward path (deriving all hidden / latent variables).

$$q_\phi(z_{1:T}|x)$$

- Extract q_ϕ and apply Jensen-inequality to get expectation over latent space.

$$\int q_\phi(z_{1:T}|x) \log \frac{p_\theta(x_0, z_{1:T})}{q_\phi(z_{1:T}|x)} dz = \mathbb{E}_{z_{1:T} \sim q_\phi(z_{1:T}|x)} \log \frac{p_\theta(x_0, z_{1:T})}{q_\phi(z_{1:T}|x)}$$

From ELBO to Loss: break down ELBO to tractable terms.

- Prior Matching: restrict latent space to Normal Distribution (**latent space**).

$$-D_{KL}(q_\phi(z_T|x_0), N(0, I)) = \int q_\phi(z_T|x_0) \log \frac{N(0, I)}{q_\phi(z_T|x_0)} dz$$

- Reconstruction: evaluate difference between generated data and real data (**data space**).

$$\mathbb{E}[\log p_\theta(x_0|z_1)] = \int q_\phi(z_1|x_0) \log p_\theta(x_0|z_1) dz$$

- Transition Quality*: only for gradual noise injection. (See Diffusion Model)

3.3 Training

The overall training loss:

$$\operatorname{argmax}_{\phi, \theta} -D_{KL}(q_\phi||p) + \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)]$$

Approximate expectation of reconstruction loss by Monte-Carlo simulation:

$$\mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] = \frac{1}{N} \sum_{i=1}^N \log p_\theta(x^{(i)}|z^{(i)}), \quad z^{(i)} \sim q_\phi(z|x^{(i)})$$

Prior matching loss is a comparison between two Gaussian distributions,

$$D_{KL}(q_\phi(z|x^{(i)})||p(z)) = \frac{1}{2}((\sigma_\phi^2(x^{(i)}))^d + \mu_\phi(x^{(i)})^T \mu_\phi(x^{(i)}) - d \log(\sigma_\phi^2(x^{(i)})))$$

Therefore, total loss can be back-propagated through encoder-decoder network.

(See [Github:CS446-MP4](#) for more details)

4 GAN

4.1 Binary Classification Game

Setup: $\eta(x) = \Pr(\text{guess}P|x)$.

$$\Pr(\text{error}|x) = \Pr(\text{guess}P, x \sim Q) + \Pr(\text{guess}Q, x \sim P) = \sum_{x \in X} \frac{1}{2} Q(x) \eta(x) + \frac{1}{2} P(x) (1 - \eta(x))$$

$$= \frac{1}{2} + \frac{1}{2} \sum_{x \in X} \eta(x) (Q(x) - P(x)) = \frac{1}{2} + \frac{1}{2} \sum_{P(x) \geq Q(x)} Q(x) - P(x) = \frac{1}{2} - \frac{1}{2} d_{TV}(P, Q)$$

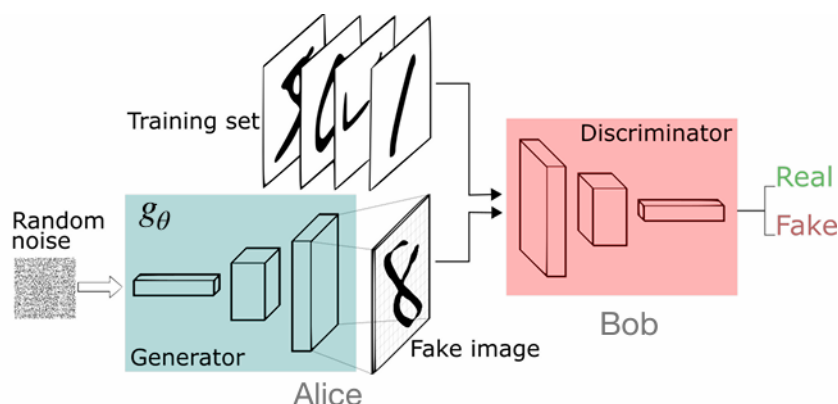
when taking optimal policy

$$\eta(x) = \begin{cases} 1, & P(x) \geq Q(x) \\ 0, & P(x) < Q(x) \end{cases}$$

Similarly, if using Cross Entropy Loss, optimal:

$$\eta(x) = \frac{P(x)}{P(x) + Q(x)}$$

4.2 Non-saturating GAN Loss (NSGAN)



$$\min_{\theta} \max_{\phi} V(g_{\theta}, f_{\phi}) = \mathbb{E}_{x \sim P_d} [\log f_{\phi}(x)] + \mathbb{E}_{z \sim P_z(z)} [\log(1 - f_{\phi}(g_{\theta}(z)))]$$

Interpretation:

- Ground Truth: True data labeled 1, generated data labeled 0.
- Generator: learn to sample from the distribution represented by the training set, try to fool Discriminator.

$$\log(1 - f_{\phi}(g_{\theta}(z))) \rightarrow 1$$

$$D^* = \arg \max_D V(G, D)$$


- Discriminator: learn to distinguish between generated and real samples.

$$\log f_{\phi}(x_d) \rightarrow 1$$

$$\log(1 - f_{\phi}(g_{\theta}(z))) \rightarrow 0$$

$$G^* = \arg \min_G V(G, D) = \arg \min_G \mathbb{E}_{z \sim P_z(z)} [\log(1 - D(G(z)))] = \arg \max_G \mathbb{E}_{z \sim P_z(z)} [\log(D(G(z)))]$$

Reason for max: larger loss focusing on low quality samples.



Limitations: training stability, behavior sensitive to hyperparameter selection. Low quality generations.

4.3 Optimization

Gradient descent-ascent algorithm (asynchronous version) with learning rate γ : Generator descent:

$$\begin{aligned}\nabla_{\theta}^{(t)} &= \nabla_{\theta} V(g_{\theta}^{(t)}, f_{\phi}^{(t)}) \\ \theta^{(t+1)} &\leftarrow \theta^{(t)} - \gamma \nabla_{\theta}^{(t)}\end{aligned}$$

Discriminator ascent:

$$\begin{aligned}\nabla_{\phi}^{(t)} &= \nabla_{\phi} V(g_{\theta}^{(t+1)}, f_{\phi}^{(t)}) \\ \phi^{(t+1)} &\leftarrow \phi^{(t)} + \gamma \nabla_{\phi}^{(t)}\end{aligned}$$

Similar for synchronous version.

(See Github:CS446-MP5.1 for more details)

4.4 Advanced Architectures

4.4.1 DCGAN

Improve discriminator architecture to reach empirically better training stability.

- No pooling, only strided convolutions.
- Leaky ReLU activations.
- Only one FC layer before softmax output.
- Batch normalization after most layers.

4.4.2 WGAN

Improve on Loss structure to get better gradients and more stable training.

- Replace sigmoid with linear activation in discriminator.
- Drop logs from objective.

$$\min_G \max_D [\mathbb{E}_{x \sim p_{data}} D(x) - \mathbb{E}_{z \sim p} D(G(z))]$$

- Clip weights to range $[-c, c]$ to ensure smoothness of discriminator.

4.4.3 LSGAN

Improve on Loss structure with least squares cost to get higher-quality images.

$$L_D = \mathbb{E}_{x \sim p_{data}} (D(x) - 1)^2 + \mathbb{E}_{z \sim p} (D(G(z)))^2$$

$$L_G = \mathbb{E}_{z \sim p} (D(G(z)) - 1)^2$$

4.4.4 ProgressiveGAN

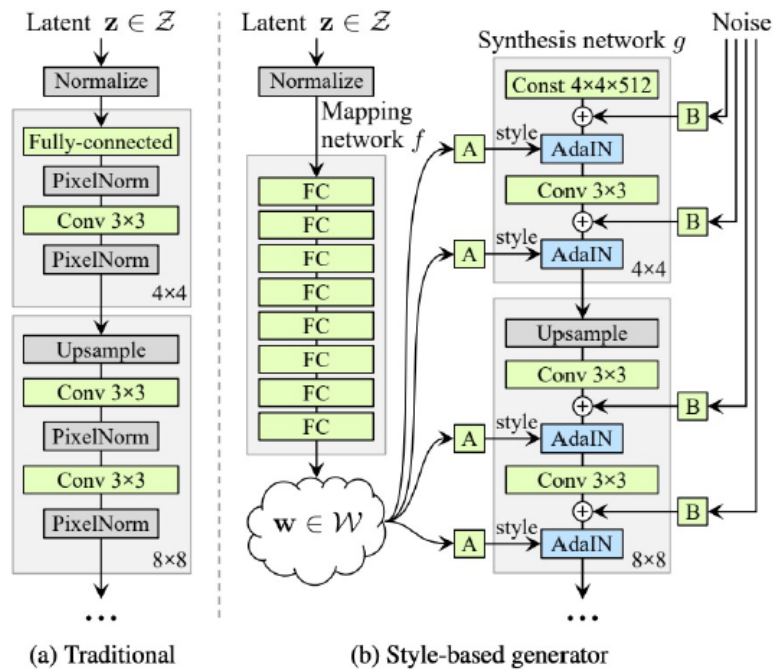
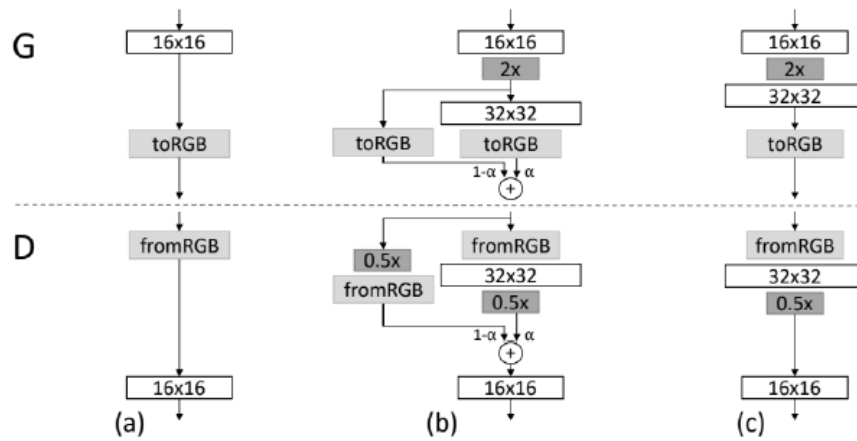
Start from training lower-resolution models, then gradually add layers corresponding to higher-resolution outputs and train.

4.4.5 StyleGAN

Improve based on ProgressiveGAN.

- Start generation with constant instead of noise vector.

Transition from 16x16 to 32x32 images



- Noise vector is transformed to latent vector w (style codes, control adaptive instance normalization or scaling and biasing of each feature map).
- Add noise after each convolution and before nonlinearity to enable stochastic details.

4.4.6 BigGAN

Scale up self-attention GAN to high resolution images.

4.4.7 StyleGAN-XL

Introduce multiple discriminators operating on projections P_l from a fixed pre-trained feature space.

$$V(G, D) = \sum_l (\mathbb{E}_{x \sim p_{data}} \log D_l(P_l(x)) + \mathbb{E}_{z \sim p} \log (1 - D_l(P_l(G(z)))))$$

where each P_l returns a feature map of a different resolution by applying random cross-channel mixing and cross-scale mixing to a pre-trained network.

Cross-channel mixing (CCM)

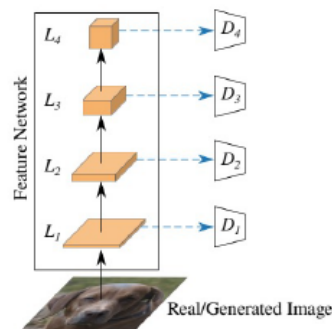


Figure 2: **CCM** (dashed blue arrows) employs 1×1 convolutions with random weights.

Cross-scale mixing (CSM)

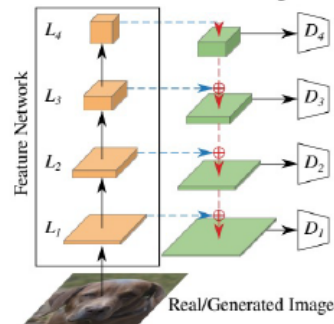


Figure 3: **CSM** (dashed red arrows) adds random 3×3 convolutions and bilinear upsampling, yielding a U-Net.

4.4.8 GigaGAN

Text-to-Image architecture.

Generator: CLIP text encoder embed text to latent code and attention values.

Discriminator: NSGAN, CLIP contrastive loss.

(See Github:CS444-MP4 for more details)

4.5 Evaluation

4.5.1 Human Studies

Turing Tests.

4.5.2 Inception Score (IS)

- Pass generated samples x through an image classifier (InceptionNet), compute posterior class distributions $P(y|x)$ and marginal distribution $P(y)$.
- Compute Inception Score, the higher the better.

$$\text{IS}(G) = \exp [\mathbb{E}_{x \sim G} \text{KL} (P(y|x) || P(y))] = \exp \left[\sum_{x \in G} P(y|x) \log \frac{P(y|x)}{P(y)} \right]$$

Higher $P(y|x)$ means higher quality that its prediction is closer to ground-truth class. Lower $P(y)$ means generated samples are diverse that contain objects of many classes. So overall the higher IS the better.

Limitation: can't detect overfitting (memorize training data) or mode-dropping (output a single image per class).

4.5.3 Frechet Inception Distance (FID)

- Pass generated samples x through an image classifier (InceptionNet), compute activations for a chosen layer.
- Estimate multivariate mean and covariance of activations, compute Frechet Inception Distance (FID) for a chosen layer. Limitation: can't detect overfitting (memorize training data) or mode-dropping (output a single image per class).

Can detect mode dropping, but can't detect overfitting.

5 Diffusion

5.1 Setup

Similar to VAE, diffusion model contains an encoder and decoder to approximate the underlying true conditional distribution.

$$\begin{aligned} q_{\phi}(x_t|x_{t-1}) &\approx p(x_t|x_{t-1}) \\ p_{\theta}(x_t|x_{t+1}) &\approx p(x_t|x_{t+1}) \end{aligned} \tag{4}$$

- Forward Path: gradually add noise till Gaussian.
- Backward Path: gradually denoise till image output.

Different from VAE, the latent variables are also x_t (same dimension as input image) instead of z .

5.2 Encoder

Manually control the noise infection process so that x_t is deterministic given x_0

$$q_\phi(x_t|x_{t-1}) = N(x_t|\sqrt{\alpha_t}x_{t-1}, (1 - \alpha)I)$$

5.2.1 Objective

$$\lim_{t \rightarrow \infty} x_t = N(0, I) \quad (5)$$

which is pure Gaussian Noise.

5.2.2 Derivation

Assume

$$x_t = ax_{t-1} + b\epsilon_t$$

where $\epsilon_t \sim N(0, I)$ is pure Gaussian Noise. Then

$$\begin{aligned} x_t &= ax_{t-1} + b\epsilon_t = a(ax_{t-2} + b\epsilon_{t-1}) + b\epsilon_t = a^2x_{t-2} + ab\epsilon_{t-1} + b\epsilon_t \\ &= \dots = a^tx_0 + b \sum_{i=0}^{t-1} a^i \epsilon_{t-i} \end{aligned} \quad (6)$$

Since ϵ is pure Gaussian Noise $N(0, I)$, $\mathbb{E}[\epsilon_{t-i}] = 0$, $Var(\epsilon_{t-i}) = I$.

Since a^tx_0 is a constant, $Var(a^tx_0 + Y) = Var(Y)$.

By linearity of Gaussian, x_t is still a Gaussian Distribution with

$$\begin{aligned} \mathbb{E}[x_t] &= \mathbb{E}[a^tx_0 + b \sum_{i=0}^{t-1} a^i \epsilon_{t-i}] = \mathbb{E}[a^tx_0] + b \sum_{i=0}^{t-1} a^i \mathbb{E}[\epsilon_{t-i}] = a^tx_0 \\ Var(x_t) &= Var(a^tx_0 + b \sum_{i=0}^{t-1} a^i \epsilon_{t-i}) = Var(b \sum_{i=0}^{t-1} a^i \epsilon_{t-i}) = b^2 \sum_{i=0}^{t-1} a^{2i} I \end{aligned}$$

To make (6) reach objective (5),

$$\begin{cases} \lim_{t \rightarrow \infty} a^t x_0 = 0 \\ \lim_{t \rightarrow \infty} b^2 \sum_{i=0}^{t-1} a^{2i} = 1 \end{cases} \rightarrow \begin{cases} |a| < 1 \\ \frac{b^2}{1-a^2} = 1 \end{cases}$$

Set $a^2 = \alpha_t$, then $a = \sqrt{\alpha_t}$, $b = \sqrt{1 - \alpha_t}$,

$$\begin{aligned} x_t &= \sqrt{\alpha_t} x_{t-1} + \sqrt{1 - \alpha_t} \epsilon = \sqrt{\alpha_t} (\sqrt{\alpha_{t-1}} x_{t-2} + \sqrt{1 - \alpha_{t-1}} \epsilon') + \sqrt{1 - \alpha_t} \epsilon \\ &= \sqrt{\alpha_t \alpha_{t-1}} x_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \epsilon = \dots = \sqrt{\prod_{i=1}^t \alpha_i} x_0 + \sqrt{1 - \prod_{i=1}^t \alpha_i} \epsilon \end{aligned} \quad (7)$$

Therefore, all x_t are tractable given x_0 .

5.3 ELBO

Objective is maximizing the log likelihood of decoder predicting real image.

$$\max_{\theta} \log p_{\theta}(x_0)$$

Since $p_{\theta}(x_0)$ itself is intractable, but conditional probability (4) is tractable by decoder, transform objective to marginalize latent variables

$$\log p_{\theta}(x_0) = \log \int p_{\theta}(x_{0:T}) dx_{1:T}$$

Since x_t itself is intractable, but the manually controlled encoding process makes it determined by encoder, make $p_{\theta}(x_{0:T})$ based on output of encoder

$$= \log \int q_{\phi}(x_{1:T}|x_0) \frac{p_{\theta}(x_{0:T})}{q_{\phi}(x_{1:T}|x_0)} dx_{1:T}$$

By Jensen-inequality:

$$\geq \int q_{\phi}(x_{1:T}|x_0) \log \frac{p_{\theta}(x_{0:T})}{q_{\phi}(x_{1:T}|x_0)} dx_{1:T} = \mathbb{E}_{q_{\phi}(x_{1:T}|x_0)} \log \left[\frac{p_{\theta}(x_{0:T})}{q_{\phi}(x_{1:T}|x_0)} \right] = ELBO$$

5.4 Disastrous Loss Function

Intuitively, given a latent state x_t , we can compute encoder ($q_\phi(x_t|x_{t-1})$) and decoder ($p_\theta(x_t|x_{t+1})$) predictions from two directions and form a loss compared with the given latent state x_t .

Since encoder and decoder are both one-step conditional prediction, by Markov Chain,

$$\begin{aligned} ELBO &= \mathbb{E}_{q_\phi(x_{1:T}|x_0)} [\log \frac{p(x_T) \prod_{t=2}^T p_\theta(x_{t-1}|x_t) p_\theta(x_0|x_1)}{q_\phi(x_T|x_{T-1}) \prod_{t=1}^{T-1} q_\phi(x_t|x_{t-1})}] \\ &= \mathbb{E}_{q_\phi(x_{1:T}|x_0)} [\log \frac{p(x_T) p_\theta(x_0|x_1)}{q_\phi(x_T|x_{T-1})}] + \mathbb{E}_{q_\phi(x_{1:T}|x_0)} [\log \prod_{t=1}^{T-1} \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_t|x_{t-1})}] \end{aligned}$$

The second term can be further simplified,

$$\mathbb{E}_{q_\phi(x_{1:T}|x_0)} [\log \prod_{t=1}^{T-1} \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_t|x_{t-1})}] = \sum_{t=1}^{T-1} \mathbb{E}_{q_\phi(x_{t-1}, x_{t+1}|x_0)} [\log \frac{p_\theta(x_t|x_{t+1})}{q_\phi(x_t|x_{t-1})}]$$

Since the distribution $q_\phi(x_{t-1}, x_{t+1}|x_0)$ is not directly tractable, this loss function is not feasible, which is disastrous.

5.5 Reverse Encoder

After a close inspection of this disastrous situation, we can see that the opposite direction of prediction is the core obstacle. Therefore, we can reverse the transition of encoder by Bayes Rule,

$$q_\phi(x_t|x_{t-1}, x_0) = \frac{q_\phi(x_{t-1}|x_t, x_0) q_\phi(x_t|x_0)}{q_\phi(x_{t-1}|x_0)} \quad (8)$$

Since encoder is manually controlled, $q_\phi(x_t|x_0)$ and $q_\phi(x_{t-1}|x_0)$ are both tractable (7), now we consider the remaining term $q_\phi(x_{t-1}|x_t, x_0)$.

$$\begin{aligned} q_\phi(x_{t-1}|x_t, x_0) &= q_\phi(x_t|x_{t-1}, x_0) \frac{q_\phi(x_{t-1}|x_0)}{q_\phi(x_t|x_0)} \\ &= C_1 \exp(-\frac{1}{2}(\frac{(x_t - \sqrt{\alpha_t}x_{t-1})^2}{1 - \alpha_t} + \frac{(x_{t-1} - \sqrt{\prod_{i=1}^{t-1} \alpha_i}x_0)^2}{1 - \prod_{i=1}^{t-1} \alpha_i} - \frac{(x_t - \sqrt{\prod_{i=1}^t \alpha_i}x_0)^2}{1 - \prod_{i=1}^t \alpha_i})) \\ &= C_1 \exp(-\frac{1}{2}(\frac{\alpha_t}{1 - \alpha_t} + \frac{1}{1 - \prod_{i=1}^{t-1} \alpha_i})x_{t-1}^2 - (\frac{2\sqrt{\alpha_t}}{1 - \alpha_t}x_t + \frac{2\sqrt{\prod_{i=1}^{t-1} \alpha_i}}{1 - \prod_{i=1}^{t-1} \alpha_i}x_0)x_{t-1} + C_2) \end{aligned}$$

where C_1 and C_2 are constants independent of x_{t-1} , so $\mu = -\frac{B}{2A}$,

$$\begin{aligned}\mu(x_t, x_0) &= \left(\frac{\sqrt{\alpha_t}}{1 - \alpha_t} x_t + \frac{\sqrt{\prod_{i=1}^{t-1} \alpha_t}}{1 - \prod_{i=1}^{t-1} \alpha_t} x_0 \right) / \left(\frac{\alpha_t}{1 - \alpha_t} + \frac{1}{1 - \prod_{i=1}^{t-1} \alpha_t} \right) \\ &= \frac{\sqrt{\alpha_t}(1 - \prod_{i=1}^{t-1} \alpha_t)}{1 - \prod_{i=1}^t \alpha_t} x_t + \frac{\sqrt{\prod_{i=1}^{t-1} \alpha_t}(1 - \alpha_t)}{1 - \prod_{i=1}^t \alpha_t} x_0\end{aligned}\quad (9)$$

From (6),

$$x_0 = \frac{1}{\sqrt{\prod_{i=1}^t \alpha_t}} \left(x_t - \sqrt{1 - \prod_{i=1}^t \alpha_t} \epsilon \right)$$

Plug in x_0 ,

$$\mu(x_t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \prod_{i=1}^t \alpha_t}} \epsilon \right)$$

Similarly, for variance,

$$\Sigma(x_t) = \frac{(1 - \alpha_t)(1 - \sqrt{\prod_{i=1}^{t-1} \alpha_t})}{1 - \prod_{i=1}^t \alpha_t} I$$

Therefore, $q_\phi(x_{t-1}|x_t, x_0)$ is also a Gaussian distribution completely rely on x_t , which has the same direction as decoder.

5.6 Loss Function

Up till now, encoder and decoder can make prediction on the same direction ($x_{t-1}|x_t$). Deriving ELBO again, this time encoder and decoder prediction should overlap each other,

$$\begin{aligned}ELBO &= \mathbb{E}_{q_\phi(x_{1:T}|x_0)} \left[\log \frac{p(x_T) \prod_{t=2}^T p_\theta(x_{t-1}|x_t) p_\theta(x_0|x_1)}{\prod_{t=2}^T q_\phi(x_t|x_{t-1}) q_\phi(x_1|x_0)} \right] \\ &= \mathbb{E}_{q_\phi(x_{1:T}|x_0)} \left[\log \frac{p(x_T) p_\theta(x_0|x_1)}{q_\phi(x_1|x_0)} \right] + \mathbb{E}_{q_\phi(x_{1:T}|x_0)} \left[\log \prod_{t=2}^T \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_t|x_{t-1})} \right]\end{aligned}\quad (10)$$

Apply Bayes (8) on product in the second term,

$$\prod_{t=2}^T \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_t|x_{t-1})} = \prod_{t=2}^T \frac{p_\theta(x_{t-1}|x_t)}{\frac{q_\phi(x_{t-1}|x_t, x_0) q_\phi(x_t|x_0)}{q_\phi(x_{t-1}|x_0)}} = \prod_{t=2}^T \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_{t-1}|x_t, x_0)} \times \prod_{t=2}^T \frac{q_\phi(x_{t-1}|x_0)}{q_\phi(x_t|x_0)}$$

After cancelling same terms on nominator and denominator,

$$\prod_{t=2}^T \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_t|x_{t-1})} = \prod_{t=2}^T \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_{t-1}|x_t, x_0)} \times \frac{q_\phi(x_1|x_0)}{q_\phi(x_T|x_0)}$$

Plug back into ELBO (10), transiting the last term,

$$\begin{aligned} ELBO &= \mathbb{E}_{q_\phi(x_{1:T}|x_0)} [\log \frac{p(x_T)p_\theta(x_0|x_1)}{q_\phi(x_1|x_0)}] + \mathbb{E}_{q_\phi(x_{1:T}|x_0)} [\log \prod_{t=2}^T \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_{t-1}|x_t, x_0)} + \log \frac{q_\phi(x_1|x_0)}{q_\phi(x_T|x_0)}] \\ &= \mathbb{E}_{q_\phi(x_{1:T}|x_0)} [\log \frac{p(x_T)p_\theta(x_0|x_1)}{q_\phi(x_1|x_0)} + \log \frac{q_\phi(x_1|x_0)}{q_\phi(x_T|x_0)}] + \mathbb{E}_{q_\phi(x_{1:T}|x_0)} [\log \prod_{t=2}^T \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_{t-1}|x_t, x_0)}] \\ &= \mathbb{E}_{q_\phi(x_{1:T}|x_0)} [\log \frac{p(x_T)p_\theta(x_0|x_1)}{q_\phi(x_T|x_0)}] + \mathbb{E}_{q_\phi(x_{1:T}|x_0)} [\log \prod_{t=2}^T \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_{t-1}|x_t, x_0)}] \\ &= \mathbb{E}_{q_\phi(x_1|x_0)} [\log p_\theta(x_0|x_1)] + \mathbb{E}_{q_\phi(x_T|x_0)} [\log \frac{p(x_T)}{q_\phi(x_T|x_0)}] + \sum_{t=2}^T \mathbb{E}_{q_\phi(x_t, x_{t-1}|x_0)} [\log \frac{p_\theta(x_{t-1}|x_t)}{q_\phi(x_{t-1}|x_t, x_0)}] \\ &= \underbrace{\mathbb{E}_{q_\phi(x_1|x_0)} [\log p_\theta(x_0|x_1)]}_{\text{reconstruction}} - \underbrace{D_{KL}(q_\phi(x_T|x_0) || p(x_T))}_{\text{prior_match}} \\ &\quad - \underbrace{\sum_{t=2}^T \mathbb{E}_{q_\phi(x_t|x_0)} [D_{KL}(q_\phi(x_{t-1}|x_t, x_0) || p_\theta(x_{t-1}|x_t))]}_{\text{consistency}} = Loss \end{aligned} \tag{11}$$

Interpretation of final Loss (11):

- Reconstruction: optimize initial block, generated image is expected to be close to real image.
- Prior Match: optimize final block, $q_\phi(x_T|x_0)$ is expected to be close to $N(0, I)$.
- Consistency: optimize middle transition blocks, generated middle latent variables is expected to be close to that generated by manually controlled encoder.

In a word, diffusion model is a decoder learning from human controlled image noising process. Under the self-supervised learning literature, encoder process is generating data $q_\phi(x_t)$ and label $q_\phi(x_{t-1})$ from original input x_0 while decoder learns from these self-generated data-label pairs in a supervised manner.

5.7 Training

5.7.1 Consistency Loss

Firstly consider the core part of the training, which is minimizing the consistency loss in Loss Function (11). Intuitively it aims at making $q_\phi(x_{t-1}|x_t, x_0)$ and $p_\theta(x_{t-1}|x_t)$ as close as possible. We have already derived (9) for $q_\phi(x_{t-1}|x_t, x_0)$ from encoder, since p_θ is a neural network to train and x_t is provided as input for it to predict $p_\theta(x_{t-1}|x_t)$, we can design it in a similar and convenient way.

$$\mu_{p_\theta}(x_t) = \frac{\sqrt{\alpha_t}(1 - \prod_{i=1}^{t-1} \alpha_t)}{1 - \prod_{i=1}^t \alpha_t} x_t + \frac{\sqrt{\prod_{i=1}^{t-1} \alpha_t(1 - \alpha_t)}}{1 - \prod_{i=1}^t \alpha_t} \hat{x}_\theta(x_t) \quad (12)$$

where $\hat{x}_\theta(x_t)$ is the original image prediction by neural network p_θ given x_t .

By MLE, we can now design a training loss that minimizes the difference between means of two Gaussian distributions (9) and (12),

$$L = \frac{1}{2\sigma_q^2(t)} \left\| \underbrace{\mu_q(x_t, x_0)}_{\text{known}} - \underbrace{\mu_{p_\theta}(x_t)}_{\text{network}} \right\|^2 = \frac{1}{2\sigma_q^2(t)} \frac{\prod_{i=1}^{t-1} \alpha_t(1 - \alpha_t)^2}{(1 - \prod_{i=1}^t \alpha_t)^2} \|\hat{x}_{0_\theta} - x_0\|^2$$

5.7.2 Reconstruction Loss

After defining a proper network for decoder, since reconstruction loss (11) also contains decoder network p_θ , we also need to plug in it.

$$\begin{aligned} \log p_\theta(x_0|x_1) &\propto -\frac{1}{2\sigma_q^2(1)} \|\mu_\theta(x_1) - x_0\|^2 \\ &= -\frac{1}{2\sigma_q^2(1)} \left\| \frac{\sqrt{\alpha_1}(1 - \prod_{i=0}^0 \alpha_t)}{1 - \prod_{i=0}^1 \alpha_t} x_1 + \frac{\sqrt{\prod_{i=0}^0 \alpha_t(1 - \alpha_1)}}{1 - \prod_{i=0}^1 \alpha_t} \hat{x}_\theta(x_1) - x_0 \right\|^2 \end{aligned}$$

Since $\alpha_0 = 1$,

$$= -\frac{1}{2\sigma_q^2(1)} \|\hat{x}_\theta(x_1) - x_0\|^2$$

5.7.3 Training Algorithm

The prior match loss is guaranteed by manual control over encoder. Combining consistency and reconstruction loss, total loss takes the form:

$$Loss = - \sum_{t=1}^T \frac{1}{2\sigma_q^2(t)} \frac{\prod_{i=1}^{t-1} \alpha_i (1 - \alpha_i)^2}{(1 - \prod_{i=1}^t \alpha_i)^2} \mathbb{E}_{q(x_t|x_0)} [\|x_{0_\theta}^\wedge - x_0\|^2] \quad (13)$$

Then the algorithm is repeating this process:

- Pick a random time stamp $t \sim [1, T]$
- Draw sample from encoder $q_\phi(x_t|x_0)$
- Decoder make prediction $\hat{x}_\theta(x_t)$
- Take gradient descent step on 13

(See Github:CS446-MP5.2 for more details)

5.8 Inference

Same as VAE, inference time only uses decoder. Since for each time stamp, given input latent variable x_t , the prediction of x_{t-1} is input latent variable to stamp $t - 1$. Therefore, the output should also be a stochastic sampling (a.k.a VAE). The sampling distribution is a Gaussian with mean predicted according to (12) and variance $\sigma_q^2(t)$.

$$x_{t-1} = \mu_{p_\theta}(x_t) + \sigma_q(t)\epsilon = \frac{\sqrt{\alpha_t}(1 - \prod_{i=1}^{t-1} \alpha_i)}{1 - \prod_{i=1}^t \alpha_i} x_t + \frac{\sqrt{\prod_{i=1}^{t-1} \alpha_i (1 - \alpha_i)}}{1 - \prod_{i=1}^t \alpha_i} \hat{x}_\theta(x_t) + \sigma_q(t)\epsilon \quad (14)$$

where $\epsilon \sim N(0, I)$. Algorithm:

- Input a white noise $x_T \sim N(0, I)$ to decoder
- Decoder make prediction $\hat{x}_\theta(x_t)$
- Sample latent variable for time stamp $t - 1$ according to 14
- Repeat until x_0

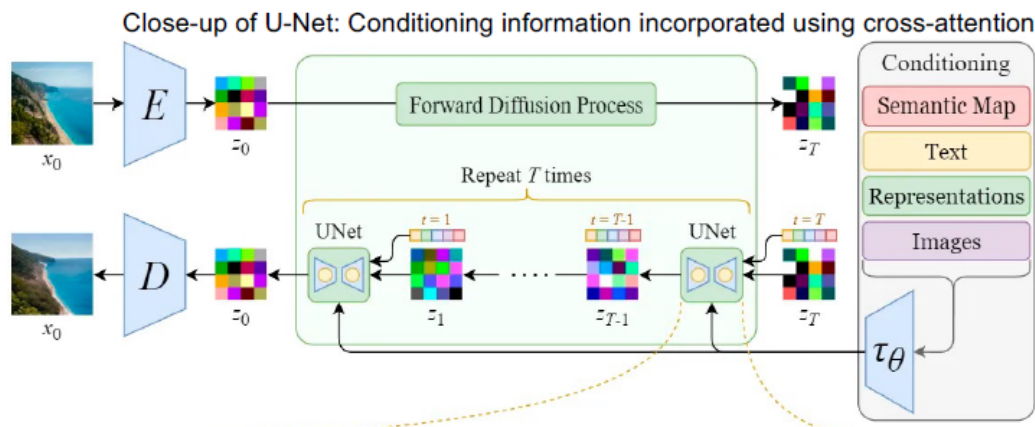
5.9 Advanced Architectures

5.9.1 DALL-E 2

Text-conditioned generation.

- CLIP text encoding: attention mechanism.
- GLIDE diffusion generator: generate image conditioned on CLIP image embedding and text prompt.

5.9.2 Latent Diffusion




Train a separate encoder and decoder to convert images to and from lower-dimensional latent space, then run conditional diffusion model in latent space instead of original size.

5.9.3 Google Imagen

A frozen LLM text encoder to embed text, a diffusion model to generate images at low resolution and upsamples to higher resolutions.

5.9.4 SDXL



Improve on pipeline, including a Base Diffusion, a Refiner Diffusion and a VAE decoder.

5.9.5 Progressive Distillation

Reducing the intermediate latent space layers progressively.

5.9.6 Latent Consistency Models

Each latent noisy level is used to predict original image and compute loss.